

An Analysis of Distributing Test Suites: Comparing Load Balancing and Sharing

Michael Camara, Colton McCurdy, Gary Miller, & Herbie Torrance

Abstract

This research focuses on utilizing distributed testing to make the execution of large test suites more efficient. The key trade-offs for taking a distributed approach to testing are identified and the impacts on using a distributed approach are analyzed. This paper will mainly focus on comparing two distributed test systems, Test Load Balancer (TLB), and our own system, Distributed Test Sharer (DTS).

Contents

1	Introduction	2
2	Background On Distributed Testing	2
3	Overview Distributed Testing Systems	4
3.1	Test Load Balancer	4
3.2	Distributed Test Sharer	5
4	Experimentation Protocol	9
4.1	Systems Specifications	9
4.2	Experimentation Protocol for TLB	10
4.2.1	Running TLB in a Distributed Fashion	10
4.3	Experimentation Protocol for DTS	11
4.3.1	Running DTS in a Distributed Fashion	11
5	Analysis of Results	12
6	Challenges	16
7	Appendix	18
7.1	TLB Data	18
7.2	DTS Data	19
7.3	Representative Source Code	20
7.3.1	Delegator	20
7.3.2	CustomServer	26
7.3.3	CustomClassLoader	29
7.3.4	TestAgent	31

1 Introduction

Testing is often considered to be one of the most important components of the software development cycle and therefore it is important to conduct a thorough and correct test suite on a software system. Many large software systems are subject to testing that can take several hours and potentially longer to execute, which can be very costly. For example, in situations where a system's test suite takes several hours or even days to execute, and one of the tests that is not executed early in the testing process fails, the cost of identifying this failure can be exceptionally high. For this reason, approaches to decreasing the cost of executing test suites is a very popular topic in software development. One of approaches to successfully reducing the cost associated with testing a system is to use distributed testing. This is an approach in which test cases are distributed or scattered across many different nodes in a network, executed locally on that given node, and the result of the test case is then returned across the network. The focus of this research is to identify, analyze, and report the results of experiments conducted on distributed testing using two different systems, Test Load Balancer (TLB) and Distributed Test Sharer (DTS). It is hypothesized that by distributing test cases and executing them in parallel on several independent nodes, the overall cost of executing a full test suite can be reduced due to the increase of computational resources and evading the limitations experienced with computation on a single node.

2 Background On Distributed Testing

In order to conduct an analysis on using a distributed approach to testing, more research on the trade-offs associated with the topic was conducted. First, it is important to understand situations in which distributed testing may be beneficial; the most obvious being when a test suite takes several hours or days to execute. It also may be necessary to conduct distributed testing if the the system of focus must be tested on different web browsers or sites, in the case of a web-based application, or even on different operating systems. In either case, it may be difficult or costly to execute these tests on a single node and therefore distributing the tests to other nodes may be beneficial.

Next, it is essential to identify both the benefits and drawbacks of testing in a distributed fashion. Some of the key benefits associated with distributed testing is that less computation is required on the CPU in which the system is executed on, which allows more work to be done locally and is potentially a large benefit in a situation where it is necessary to execute other processes alongside the test suite. More importantly, the goal is to reduce the runtime of the test suite by executing several test cases in parallel and it is an exceptionally more feasible task when using multiple CPUs. One of the most common ways in which the increase in performance is

achieved is by diminishing the cost of dependencies. Tests cases that are dependent on another case often have larger costs because their result or even execution may not occur until another case is completed. By using distributed testing, these dependent test cases can be isolated to a particular node that is responsible for executing the group of dependent cases while other nodes can execute other independent cases that otherwise would have been waiting in a queue if centralized testing was used. This situation is a clear example in which a distributed approach would be more beneficial than a centralized approach.

Though it is clear the distributed testing may be beneficial to use, there certainly are drawbacks and challenges that exist with this approach. The greatest of these drawbacks is the cost of communication required to transfer data over a network from one node to another. In order to execute a test case at a remote location, either the code to execute the test must exist at the remote location or the code must be sent across the network. This is a potentially costly situation if the required code is large and takes up a lot of memory. The situation where the code already exists at the remote location is often unlikely and therefore the test code usually must be transferred to the node that will execute it, therefore it must be determined whether communication costs of distributed testing outweighs the computational and time costs associated with executing the tests locally. Another drawback associated with distributed testing is that it is dependent on a network, which in some cases may be unavailable. If a system uses a distributed test suite and it is unable to gain access to a network, then there is no way to run the tests until a network becomes available. Another unfavorable situation may occur if there are fewer remote nodes available than required. In a situation such as this, both the communications costs and computational limitations may be experienced because all of the data is transferred to a number of locations that cannot execute the tests in a time that justifies distributing the tests.

Alongside the drawbacks of distributing testing, several challenges exist with this approach as well. One of the key challenges associated with distributed testing is the concept of load sharing versus load balancing. Load sharing occurs when the tests or load are sent to several nodes to be executed without any restrictions on how the load is allocated. Load balancing is a more complex approach in which the allocation of the load is sent to nodes such that each node is doing an equivalent amount of work. This allows the load to be executed more efficiently because it prevents situations in which a particular node is doing a significantly greater amount of work than others, or a particular node is doing little to no work at all. It is quite evident that load balancer is a more favorable approach to take, but the complexity associated with it is much larger. For example, consider a situation in which a test case is dependent on another, in this case a load balancing system must not only consider the amount of load to allocate to each node, but also when certain components of the load need to be allocated.

3 Overview Distributed Testing Systems

3.1 Test Load Balancer

Though distributed testing is a popular and rapidly developing topic in the software industry, not many systems that use this approach exist and work correctly. This is due to the fact that it is not only a relatively new concept, but the complexity of building a system that can run tests on a system and efficiently distribute them to remote nodes is exceptionally high. One existing system that successfully does this is named Test Load Balancer (TLB). TLB claims to support testing of every language on every platform and partitions the tests into subsets that can be executed in parallel on different physical or virtual machines. These subsets are executed in such a way that they all start at the same time and finish at almost the same time as well, therefore the overall time it takes to execute all of the tests can be divided by the number of test subsets that are generated. TLB assures that a particular node will execute only a number of tests proportional to the total number of tests divided by the number of available nodes and guarantees mutual exclusion and collective exhaustion of test execution, meaning that no test will be executed more than once and every test will be executed.

The TLB system is comprised of two main components, the first of which is a server that is responsible for storing test data and the second is the balancer that partitions and orders the execution of tests. By storing test data, the TLB system will reorder the execution of tests with each run based on previous results. This is an extremely beneficial feature because tests that are known to fail will be executed earlier to avoid unnecessary downtime time resulting from failing tests being executed later. Though it is very important to note that this is only achievable by using tests that are independent of one another which is not always the case. Unfortunately, large systems very often have test cases that are dependent on other tests and systems such as these are usually examples that would benefit from using distributed testing opposed to centralized testing. Also, this method can not guarantee the order in which tests are executed due to its fail early approach to executing test, therefore tests must not only be independent of other tests but also independent of the order in which they are executed.

In order to run TLB properly, it was necessary to utilize another system called Docker. This is a system that is used to effectively "pack up" a software system and all of its dependencies, such that it can be ran in the same way in any different environment. This was necessary, because not only was the software that was being tested executed in several different environments on a few different operating systems, but also TLB doesn't precisely have a means of distributing the tests across the network by itself. For this reason, another system must be used to do the actual execution of the tests, while TLB is responsible for deciding where and when to execute an individual test case. The remainder of this report will now focus on analyzing the two main

systems of focus, DTS and TLB, in more detail, and analyze the results of empirical analysis conducted on distributed testing through both load sharing using the DTS system and load balancing using the TLB system.

3.2 Distributed Test Sharer

In order to understand and further evaluate the trade-offs associated with distributed testing, we developed a novel system named the Distributed Test Sharer (DTS). This system was created using the Java programming language and consists of two main classes: **Delegator** and **CustomServer**. The relationship between these classes is illustrated in Figure 1, with communication facilitated by Java Remote Method Invocation (RMI) and use of an FTP server. The following section will outline the steps performed by DTS to distribute JUnit tests located on the **Delegator** node to any number of **CustomServer** nodes, which run the tests and return the results obtained.

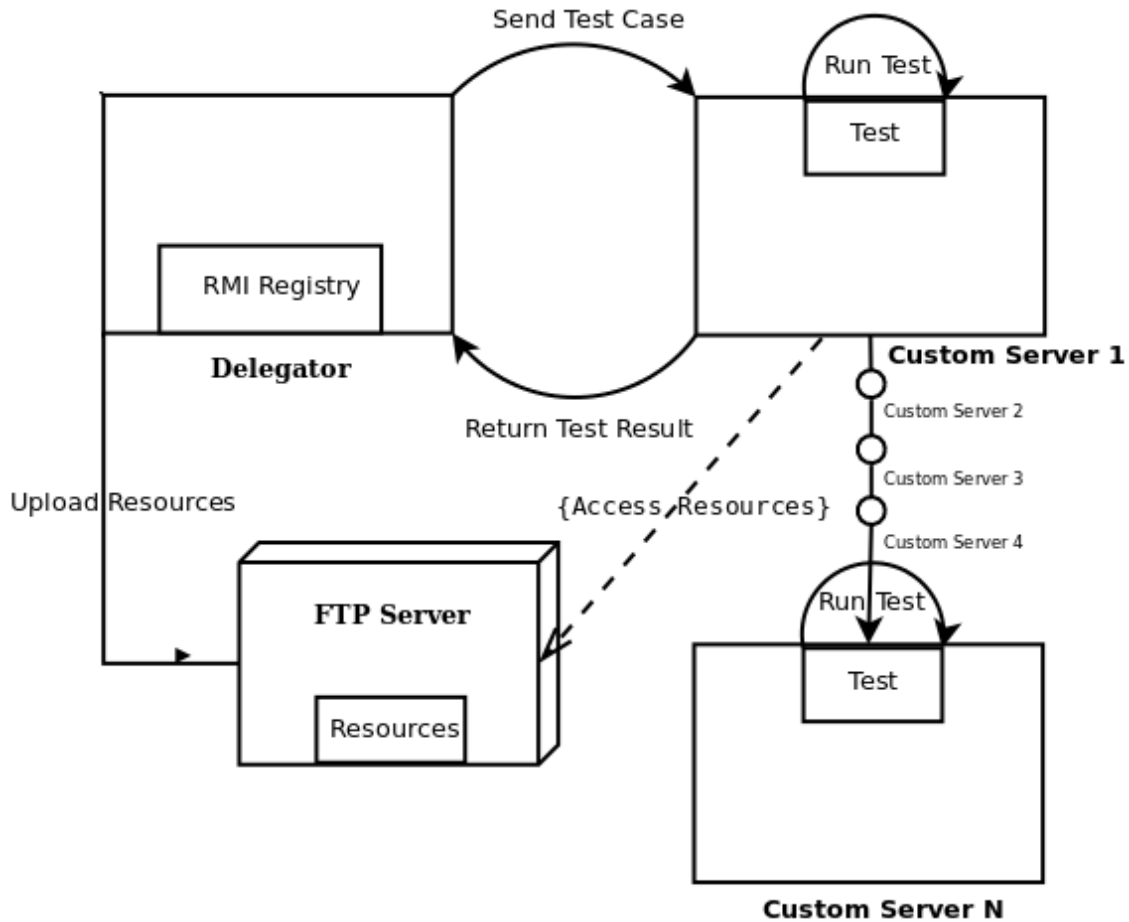


Figure 1: Architecture of the Distributed Test Sharer (DTS).

Before starting the system, the user is expected to furnish all of the JUnit test cases and any

associated resources needed to run those test cases. All files will be placed in the **ftpserver** directory, or one of its subdirectories. The top level of this directory can be filled with any non-class files or directories that might be required for test suite execution. For instance, the SchemaAnalyst test suite uses a **config** directory that contains various text files filled with different properties. In this example, the user would place the entire **config** directory immediately inside the **ftpserver** directory. Further nested in the **ftpserver** directory is a **resources** directory, which contains **bin**, **lib**, **test_singles** and **test_suites** subdirectories. For simplicity, the user should put the entirety of the directory containing their compiled classes in the **bin** directory. Continuing with the SchemaAnalyst example, this would involve placing the **org**, **paper**, and **parsedcasestudy** directories into the **ftpserver/resources/bin** directory, which are usually located in that system's **build** directory. All **jar** files used by the desired tests should be placed in the **lib** directory. Finally, any individual JUnit test cases can be put in the **ftpserver/resources/test_singles** directory, while any test suites (those that reference multiple JUnit test cases, like an **AllTests** class) can be put in the **ftpserver/resources/test_suites** directory.

DTS execution begins on the side of the **Delegator** with a few preliminary steps. First, several system properties are programmatically altered. These include **java.security.policy**, which is set to a custom policy that grants all permissions to any nodes that attempt to connect to the **Delegator** through Java RMI. Although a more stringent policy would generally be recommended to improve the security of the system, this was deemed appropriate for the purpose of prototyping the system with fewer obstacles. Further, the **java.rmi.server.hostname** property is set to the internet protocol (IP) address of the **Delegator**, which is retrieved through the command line interface at program execution. This specifies the given IP address for use with any registry that is subsequently created, allowing remote nodes to find it.

Next, a Java RMI registry is created using a port number that is currently hard coded as 12345. This port number was chosen because it resides in the range of acceptable ports for the machines used for experimentation, and generally is not assigned to any other system processes. A new instance of **Delegator** is then created and linked to the registry through use of the **rebind()** method, binding the object with a key of "Delegator." This will allow **CustomServer** nodes to access some of the methods offered by **Delegator** in future steps.

Both the **Delegator** and **CustomServer** classes extend **UnicastRemoteObject**, which allows them to be bound to and retrieved from a Java RMI registry. It facilitates the dynamic creation of stubs for such objects, obviating the static method used in previous versions of Java. Each class also implements an interface that extends **Remote: DelegatorInterface** and **CustomServerInterface**. The **Remote** marker interface is further required to successfully bind to the registry, and it allows for type checking of remote object invocations at compile time.

After the registry is created, an FTP server is created using the Apache FtpServer framework. This framework was chosen because it allows programmatic creation of an FTP server using Java, instead of relying on a third-party application that would need to be run and configured separately. Various settings are adjusted during this step, such as establishing a user name and password, setting the port and host address, and enabling concurrent logins. Most importantly, the home directory for the FTP server is set to the `ftpRootDir` variable, which is currently hard coded to the local `ftpserver` directory where the test cases and resources were previously placed. Any users that connect successfully with the server will then have immediate access to this directory and any nested directories contained within it.

The final preliminary step taken by `Delegator` is the creation of a list of JUnit test cases. Both the `ftpserver/resources/test_singles` and `ftpserver/resources/test_suites` are traversed recursively to locate all nested `class` files, which are saved as `File` objects and stored into separate global `LinkedLists`: `testSingleList` and `testSuiteList`.

After these steps are completed by the `Delegator`, it waits until the user enters the carriage return character to continue. During this wait, remote nodes can connect to the `Delegator` by running the `CustomServer` class. Importantly, two command line arguments must be specified at execution: the IP address of the `Delegator`, followed by the server's own IP address. Ideally, multicasting would be used by the `Delegator` to broadcast its address, allowing the `CustomServers` to locate it automatically. While this might be pursued in the future, the current implementation requires the `Delegator's` IP address to be known.

Once started, the `CustomServer` sets the system RMI policy in the same manner as done by the `Delegator`. Additionally, it will set the `java.system.class.loader` system property to a `CustomClassLoader`. This class extends a `URLClassLoader` and allows access to some of the normally protected methods through use of reflection. For now, it simply sets the classpath to the local `bin` and `lib` directories. Next, the `CustomServer` will locate the registry on the `Delegator`, using the IP address previously specified, and obtain a remote reference to the bound `Delegator` object using the `Naming.lookup()` method. Using this reference, it will call the `rebindServer()` method to bind the `CustomServer` to the registry. This is repeated for each `CustomServer` node until they are all bound to the registry.

After all `CustomServers` are created and connected, then the user will enter the carriage return character on the `Delegator` node to continue execution. The `Delegator` will first call the remote `updateClassLoader()` method for each of the `CustomServers`. This will add all of the relevant `class` directories and `jar` files to the classpath of each `CustomServer`, accessible via FTP protocol and the FTP server previously created. Additionally, all of the non-class files previously placed in the `ftpserver` directory on the `Delegator` node are actively retrieved through communication with the FTP server and stored locally on each server node. Although

efforts were made to obviate the need to store any files on the server, this was deemed a critical step to ensure that subsequent tests could be run successfully.

Next, the **Delegator** creates **TestAgent** objects for each **CustomServer**. These agents extend the **Thread** class and implement **Runnable**, allowing multiple agents to run their respective methods concurrently on separate threads. The **run()** method will simply do a remote call to the **runTest()** method for its assigned **CustomServer**, sending a test case and storing the **Result** generated by the server into the same **ConcurrentLinkedQueue**. These agents are further placed into another **ConcurrentLinkedQueue** for subsequent access by the **Delegator**.

Once these agents are created, the **Delegator** will sequentially iterate first through the **testSingleList**, and then through the **testSuiteList**. With the single test cases, each **File** in the list is assigned as-is to an agent in the queue, provided that agent is not currently active. The **run()** method of the agent is then started, sending the test to the assigned **CustomServer**. If the agent is active, then it is placed back into the queue and the next agent is checked, continuing in that manner. This load sharing approach ensures that each **TestAgent**, and therefore each **CustomServer**, is always running a single test case at all times. This process is similarly repeated while iterating through the test suites, with some additional steps. The test suite **File** is first converted into a byte array, which is then converted into a **Class** object using a **SimpleClassLoader**. This class loader simply overrides the **defineClass()** method of a generic **ClassLoader**, allowing this conversion to take place. Next, the test suite is further converted into an array of **Class** objects by retrieving all test cases listed under the **SuiteClasses** annotation in the suite. This array, containing all of the individual test cases in the suite, is iterated through in the same manner described previously, performing load sharing using the previously created agents.

On the **CustomServer**, there are two **runTest()** methods: one has a **File** parameter, while one has a **Class** parameter. The method for **Class** objects simply executes the static **JUnitCore.runClasses()** method on that object, which returns a **Result** object. This **Result** object contains all of the details about the successes and failures from running the test, and is returned to the caller. The **runTest()** method for **File** objects first converts the **File** to bytes, and then to a **Class** object, finally using the same **JUnitCore.runClasses()** method to run the test. This is performed by the **CustomServer** to lessen the computation required on the **Delegator** as much as possible. In both cases, running the tests is only possible due to the previously implemented **CustomClassLoader**. This allows the JVM to automatically retrieve any needed **class** or **jar** files via communication with the FTP server on the **Delegator** without any prompting from the user, and without needing to copy those files to the **CustomServer** node.

The **Delegator** continues to perform load sharing, assigning test cases to **TestAgents**. The agents continue to store the **Result** objects returned from the **runTest()** method of each

`CustomServer` into the same `ConcurrentLinkedQueue`. Once all tests have been assigned, and all `Results` have been retrieved, then the each `Result` is parsed and the total number of successes and failures are displayed for the user. Each failure is further elaborated to show why a test case might have failed. This completes the process used by the Distributed Test Sharer, with all tests located on the `Delegator` having been executed by separate, remote nodes.

4 Experimentation Protocol

4.1 Systems Specifications

- **Colton's 13-inch, Mid 2012 MacBook Pro**
 - OS / Version: OS X El Capitan version 10.11.4
 - Processor: 2.9 GHz Intel Core i7
 - Memory: 8 GB 1600 MHz DDR3
 - Startup Disk: Samsung 850 Pro SSD
- **Herbie's Hp Envy TS 17 Notebook PC**
 - OS / Version: 64-bit Windows 10
 - Processor: 2.4 GHz Intel Core i7-4700MQ
 - Memory: 16 GB
 - Startup Disk: Samsung 850 Pro SSD
- **Herbie's HP Dv7-6c95dx Notebook PC**
 - OS / Version: 64-bit Windows 10
 - Processor: 2.20 GHz 2nd generation Intel Core i7-2670QM
 - Memory: 8 GB
 - Startup Disk: Samsung 280 Pro SSD
- **Herbie's Fujitsu lifebook t2020 (x2)**
 - OS / Version: 32-bit Ubuntu 14.04
 - Processor: 1.2 GHz Intel Core 2 Duo U9400
 - Memory: 3 GB
- **Michael's Dell Inspiron**
 - OS / Version: 64-bit Windows 8.1
 - Processor: 2.16 GHz Intel Core Duo Celeron n2830
 - Memory: 4 GB

4.2 Experimentation Protocol for TLB

For TLB we ran a Docker container with Java SE Development Kit 8u91 and Apache Ant version 1.9.4. These were both necessary for running the existing tool in the Test Load Balancer (TLB). TLB is compatible with JDK versions as old as version 6, but to keep our experiments as consistent as possible, we installed the latest version (8u91) on all of the machines being used for testing. Having the newest version of Apache Ant was not entirely necessary, but again, we wanted to have the most recent versions of software for this experiment due to possible optimizations in the newest versions.

4.2.1 Running TLB in a Distributed Fashion

We ran the TLB tool five times for all configurations (e.g. three nodes on the schema-analyst test suite). We decided to only run each configuration five times due to time constraints. We had it calculated that if we ran each configuration for the thirty trials suggested by Traeger et al. that our experiments combined would take approximately 20-hours to run, requiring us to sit at our computers and monitor the tests constantly for this amount of time [1].

Also, for the TLB experimentation setup, we ran Docker containers on each node. There were three total nodes, all of which were the available personal laptops of a team member. The system specifications for the machines used in this experiment can be found in section 4.1. We would have liked to run the experiments in a more controlled and consistent environment by running the Docker containers with TLB on the Alden lab machines, but were denied permission to install Docker for various potential security reasons. To read more about the challenges faced during this assignment, read section 6.

The way TLB is designed requires a user to start the **TLBServer** on one node and start the clients with the **TLBBalancer** on other nodes. Therefore, to collect the data necessary, we started the server on the 13-inch, Mid 2012 MacBook Pro and ran the Docker containers on two other machines as well as on the MacBook Pro. This was the setup for running three clients with the TLB system. For two and a single client, we removed Michael's Dell Inspiron, which performed the worst when running the test suites in the Docker container. Specifically, for two nodes, we used Herbie's HP Envy as one client node and Colton's MacBook Pro as the other client as well as the server. Ideally, we would have liked to only run the server on Colton's MacBook Pro and the clients on their own individual machines, but again, Michael's laptop took substantially longer time to execute its portion of the test suite. However, for a single node, we did not have to worry about running a server on a different machine. We ran the Apache Ant task `test` which for SchemaAnalyst's test suite called five iterations of the test suite and for Apache Ivy ran the test suite provided with the system.

It was the case that in order to make test suite execution time long enough to see substantial

differences between number of nodes, we needed to run multiple iterations of the same test to artificially increase test suite execution time. Additionally, for comparison to a test-sharing system we needed to run this modified version of SchemaAnalyst’s test suite.

This is the protocol we followed throughout the experimentation phase of this final laboratory assignment. We acknowledge that there were a lot of places where we did not apply the industry-standard protocols for testing, however due to time constraints and availability of machines, this is the best protocol that we could follow to ensure that the data that we collect was as accurate and consistent as possible.

4.3 Experimentation Protocol for DTS

Similar to running TLB, there were many steps that were taken to ensure that the quality of experimentation for DTS was the highest possible. While running tests for a distributed test suite it is important to consider the fact that outside factors may have an impact on the results. Since this is the case, we listed all possible factors which could hinder a smooth analysis of the program. We then used our knowledge of these potential problems to minimize their impact.

4.3.1 Running DTS in a Distributed Fashion

The first step in our experimentation was choosing nodes which could host our system. The computer system implemented in Alden Hall at Allegheny College was unable to run our customized system. Since this was the case we chose to gather personal computers from different group members for testing. The difference in testing between our DTS and TLB is that some computers which were viable for DTS were not viable for TLB. The computers used for DTS testing were Herbie’s HP Envy, Herbie’s DV7, Herbie’s Fujitsu Lifebook X2, and Michael’s Dell Inspiron. Since the format of our implementation has a delegator, we set aside Herbie’s HP Envy solely for this purpose. This means that we had a total of 4 nodes to start running the tests on.

As mentioned previously, these computers ran different operating systems. This means it was necessary to configure the operating systems separately. The main concern was that automatic updates and background processes were double checked to be inactive during testing. After this was done, we were able to run the tests for 5 trial per number of nodes. To clarify, we ran 5 good trials for each 4,3,2,and 1 node. While doing this it is important to note that the network used for this experiment was public. Since the network was public we had to ensure that the network was not being slowed by other users. In order to combat this situation we ran the tests on the network around 5 A.M Eastern Time on a Saturday. This would most likely be the least active time for the network on a college campus. Each test from this point was run successfully while considering a few different things.

The first thing to consider was that data may be cached after running the program multiple

times. In order to combat this we ran the program 5 times before collecting any data. Another concern was that Java’s garbage collector would interfere. This being said, we discarded any data was more than 8 percent deviation from the average since it is very difficult to infer what data is being skewed by the garbage collector. The tests run on 4 nodes was run with all computers. The tests run on 3 nodes was run with Herbie’s DV7 and T2020 X2 laptops. The tests run on 2 nodes was run with the T2020 lifebooks only. This was reduced to a single T2020 lifebook for the single node.

Similar to testing the TLB system, there were places where we did not follow industry standards of testing. We acknowledge that there may be some cases that need further explored. However, as talked about above, we did try to combat this at every possible instance. Overall, experimentation with the DTS system was a success and the results of these experiments are further analysed in the next section,

5 Analysis of Results

The first analysis will be of running Apache Ivy’s test suite on the Test Load Balancer tool. On average, on one node, running the Apache Ivy test suite took approximately six minutes to execute. To clarify, when I refer to running on a single node, I am referring to running the test suite using the `ant test` task in the `build.xml` instead of the `ant test.load.balanced` task.

Where we expected to see a significant performance increase—especially with the Apache Ivy test suite— was when adding additional nodes. However, we did not see this dramatic increase in performance and actually experienced a decrease in performance when adding the third node. I will note that the third node that was added was Michael’s Dv7-6c95dx which we noticed took substantially longer to execute test suites than other nodes on average.

Therefore, we see the greatest performance when balancing the work using the Test Load Balancer tool to two nodes. When executing the Apache Ivy test suite on two remote nodes, we notice that each node is able to complete their portion of the test in approximately 1.25 minutes. Refer to Figures 2 and 3.

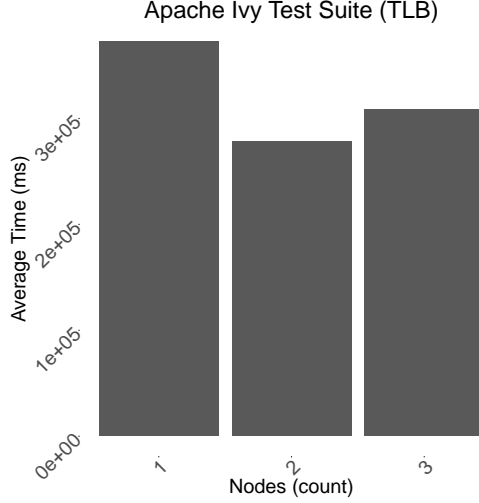


Figure 2: A bar chart displaying average time in milliseconds to run Apache Ivy’s test suite on TLB




Figure 3: A boxplot displaying average time in milliseconds to run Apache Ivy’s test suite on TLB

Secondly, and more importantly, we wanted to analyze the effectiveness of the Test Load Balancer tool at balancing SchemaAnalyst’s test suite. This is more important to use in particular because the Distributed Test Sharing tool is—in its current state—only able to execute this test suite. With that said, we are able to compare direct performances of the tools. However, again, there are some inconsistencies such as which laptops were available for the systems to be run on.

Similarly to the Apache Ivy’s test suite for the Test Load Balancing tool, the TLB tool experiences its best performance with two nodes. We hypothesize that this is for the same reason as noted previously regarding to the performance of the one node. However, here when the third node is introduced, running the test suite locally, one a single node is more performant. This makes it not worth distributing tests to more than two nodes.

SchemaAnalyst’s test suite on a single node took on average 15 seconds to execute. We do note that there are a lot of failing tests and missing resources. This is why we hypothesize that this execution time is so minimal. A lot of the dependencies necessary for this test suite are database management systems that were not installed on many of the machines used in this experiment.

In conclusion, for the TLB tool, it is not worth running either Apache Ivy’s or SchemaAnalyst’s test suite on more than two nodes. For both of these test suites, it is worth balancing the work to two remote nodes using the TLB tool communicating with the TLBServer to POST, PUT and GET historical test data for a given test suite.

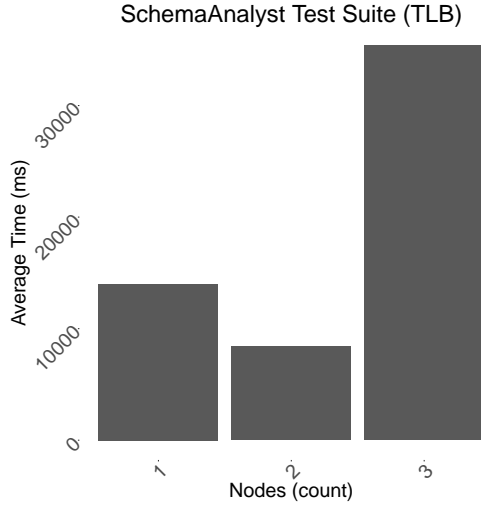


Figure 4: A bar chart displaying average time in milliseconds to run SchemaAnalyst’s test suite on TLB




Figure 5: A boxplot displaying average time in milliseconds to run SchemaAnalyst’s test suite on TLB

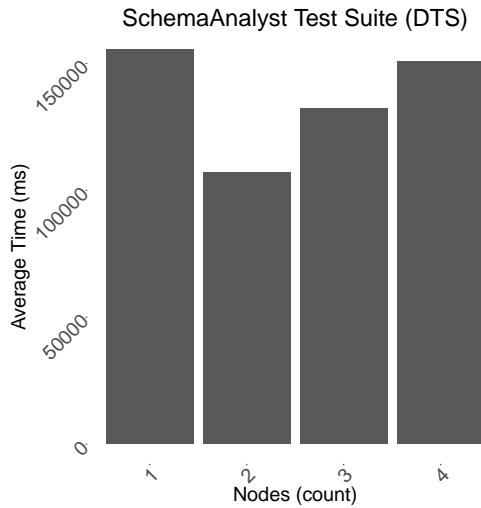


Figure 6: A bar chart displaying average time in milliseconds to run SchemaAnalyst’s test suite on DTS




Figure 7: A boxplot displaying average time in milliseconds to run SchemaAnalyst’s test suite on DTS

Finally, we are interested in the performance of the system created by Michael, Distributed Test Sharer (DTS). As previously noted, we were only able to run SchemaAnalyst’s test suite on DTS.

First, in Figures 6 and 7, we see that running the test suite on a single node takes about 2.5 minutes to execute. Then, we notice that when sharing the work with another node, we see a pretty substantial performance increase, reducing the overall time to execute the test suite to a

little over a minute and a half. This is a total reduction of about a whole minute.

However, we see a trend when observing these visualizations. We notice that as we add more nodes to share work with there is a direct correlation to the increase in time to execute the test suite. This is due to the overwhelming communication costs between these larger number of nodes. We are able to eliminate the costs of actually transferring the files necessary for running the test suite by accessing them through a file transfer protocol. Had we added this additional cost on top of already expensive communication, running the DTS tools as a sharing mechanism for SchemaAnalyst’s test suite would make it less worth much more quickly.




Figure 8: A direct performance comparison of the TLB and DTS tools for running SchemaAnalyst’s test suite

In Figure 8, we notice that both TLB and DTS behave similarly when distributing the SchemaAnalyst test suite. Running the test suite on a single node takes longer than the most performant number of nodes—that being two total nodes. At two nodes, both systems perform their absolute best and only decrease in overall performance from there.

Remember, we were limited to only comparing the test suite of the SchemaAnalyst tool and on very few nodes due to the lack of permissions and available resources. We would have liked to run these two systems on many more similar nodes to alleviate any inconsistencies, however this was just not possible for the scope of this final laboratory assignment.

6 Challenges

Various challenges were encountered while developing the Distributed Test Sharer, with many of these challenges revolving around the use of Java RMI for communication between nodes. There were a variety of different settings and properties that I didn't originally know about, despite using this technique in a previous assignment. These included either setting the RMI host name and security policy through the `-Djava` command line parameter, or using the `System.setProperty()` method programmatically. Omitting these steps produced a variety of different, confusing errors, and required much troubleshooting to figure out how to fix them. There were also problems getting the RMI to work on the computers in Alden Hall, which ultimately prevented those computers from being used in our experiments. Originally, these problems arose from the fact that the Java RMI registry would occasionally use a random port number instead of the one specified in the source code. This created much confusion originally, since the Java API includes various method signatures that require a port number parameter; yet it seemed as though such specification did not have the expected effect. Ultimately, using a custom `SocketServerFactory` and adjusting the parameters for the `UnicastRemoteObject` constructor fixed those initial problems involving randomly assigned ports. However, new problems arose with remote objects being removed from the registry after binding, perhaps due to some garbage collection issue. These problems were not present during testing of DTS on personally owned computers, and unfortunately they could not be resolved despite significant time investment.

Another challenge with DTS was getting the tests to run on a separate node without error, and specifically without needing to copy over any files to the servers. I used a custom `URLClassLoader` to originally overcome this challenge, which is a vital component to the system. Although this allows the JVM to access the FTP server for `class` and `jar` file, it doesn't redirect the paths used for other kinds of files. For instance, the SchemaAnalyst test suite makes use of a `config` directory that contains various properties files. Although I could specify the exact FTP URL for the class loader to locate this directory, the test cases would not find it when they ran on the server. I tried to adjust the system property for the current working directory (`user.dir`), among other fixes, but could never bypass this error completely. Ultimately, in order to get the tests to work as expected, I need to add a step to the `updateClassLoader()` method on the `CustomServers` to retrieve and store some files from the FTP server. This was unfortunate, since a notable benefit of this system is the lack of resource overhead for the servers; but it was necessary in order to have a functional system within our time constraints. This general issue with file dependencies further complicated running DTS on complex systems like `ant-ivy` and `tomcat`, despite the accommodations made to DTS. This prevented DTS to be run on similar systems as TLB, and it would be one of the first areas of improvement if development on this system were to continue.

In addition to the challenges Michael encountered while implementing the DTS tool, there were similar challenges faced while configuring, analyzing and running the existing, Test Load Balancing tool.

First, while the system is on GitHub (<https://github.com/test-load-balancer/tlb>), it is not a working version. GitHub is initially where I found the tool, but in the tool repository the creators direct you to the tool website (<http://test-load-balancer.github.io/>). The tool website is filled with data and comments about the tool, however when it came time to configuring and running the tool, these processes were not made absolutely clear. No where on this page does it say they do not support actually running the clients in a distributed fashion.

Also, while there is a configuration page containing documentation for all of the necessary environment variables, the provided examples and descriptions are not clear at describing all of the variables. Luckily, we were able to reach out and hear back from the creators of this tool. They were able to guide us in the correct direction. But even without the help of the authors, we were able to find the answer to our question regarding how to configure the distribution of a test suite using the TLB tool was answered in a blog post by one of the creators of the TLB tool. Even here though, there were only one or two sentences on the topic. Maybe it was made clear somewhere else, but to me, this fact was not evident.

In order to run the TLB tool in a distributed fashion, we were required to use multiple nodes. We were able to accomplish this by creating containers—which are virtually easy-to-configure virtual machines with all of the necessary dependencies for running TLB. While we did run into one issue with Docker, that being installing Docker on a 32-bit machines, overall, Docker was a pretty easy-to-use system, especially for Mac OS X.

As mentioned, the only issue we had with Docker directly was when trying to install it on an unsupported build of Windows. To combat this compatibility issue, we asked permission to have Docker installed on the Alden lab machines. If Docker were installed on the lab machines, our experiments would have been much more interesting, consistent and more detailed.

The final issue that I encountered was when trying to run experiments. I knew the TLB tool had been working the evening before trying to run the experiments and I had not manipulated any of the files and the TLB system was not communicating or running the test suites. I tried a number of ports to no avail. Finally, I decided to try the system at home once again and everything worked flawlessly, therefore requiring me to haul all of the necessary machines for the experiment to my house.

7 Appendix

7.1 TLB Data

```
tool,number_nodes,trial,time
ant-ivy,3,1,381
ant-ivy,3,2,288
ant-ivy,3,3,284
ant-ivy,3,4,285
ant-ivy,3,5,303
ant-ivy,2,1,270
ant-ivy,2,2,287
ant-ivy,2,3,275
ant-ivy,2,4,277
ant-ivy,2,5,281
ant-ivy,1,1,338
ant-ivy,1,2,399
ant-ivy,1,3,380
ant-ivy,1,4,372
ant-ivy,1,5,375
schema-analyst,3,1,34
schema-analyst,3,2,34
schema-analyst,3,3,40
schema-analyst,3,4,34
schema-analyst,3,5,35
schema-analyst,2,1,8
schema-analyst,2,2,8
schema-analyst,2,3,10
schema-analyst,2,4,8
schema-analyst,2,5,8
schema-analyst,1,1,14
schema-analyst,1,2,13
schema-analyst,1,3,16
schema-analyst,1,4,13
schema-analyst,1,5,14
```

7.2 DTS Data

```
tool,number_nodes,trial,time
schema-analyst,4,1,153755
schema-analyst,4,2,138035
schema-analyst,4,3,150007
schema-analyst,4,4,159435
schema-analyst,4,5,152664
schema-analyst,3,1,143639
schema-analyst,3,2,132070
schema-analyst,3,3,139330
schema-analyst,3,4,131438
schema-analyst,3,5,114332
schema-analyst,2,1,109954
schema-analyst,2,2,105859
schema-analyst,2,3,107833
schema-analyst,2,4,107678
schema-analyst,2,5,104399
schema-analyst,1,1,134865
schema-analyst,1,2,142017
schema-analyst,1,3,186619
schema-analyst,1,4,172291
schema-analyst,1,5,142197
```

7.3 Representative Source Code

7.3.1 Delegator

```
1 package delegator;
3 import java.io.File;
import java.io.FileInputStream;
5 import java.rmi.RemoteException;
import java.rmi.registry.LocateRegistry;
7 import java.rmi.registry.Registry;
import java.rmi.server.UnicastRemoteObject;
9 import java.util.ArrayList;
import java.util.Iterator;
11 import java.util.LinkedList;
import java.util.Scanner;
13 import java.util.concurrent.ConcurrentLinkedQueue;

15 import org.apache.commons.io.FilenameUtils;
import org.apache.ftpserver.ConnectionConfigFactory;
17 import org.apache.ftpserver.DataConnectionConfigurationFactory;
import org.apache.ftpserver.FtpServer;
19 import org.apache.ftpserver.FtpServerFactory;
import org.apache.ftpserver.filesystem.nativefs.NativeFileSystemFactory;
21 import org.apache.ftpserver.ftplet.Authority;
import org.apache.ftpserver.ftplet.FtpException;
23 import org.apache.ftpserver.ftplet.UserManager;
import org.apache.ftpserver.listener.ListenerFactory;
25 import org.apache.ftpserver.usermanager.impl.BaseUser;
import org.apache.ftpserver.usermanager.impl.WritePermission;
27 import org.apache.log4j.Logger;
import org.apache.log4j.varia.NullAppender;
29 import org.junit.runner.Result;
import org.junit.runner.notification.Failure;
31 import org.junit.runners.Suite.SuiteClasses;

33 import server.CustomServerInterface;

35 public class Delegator extends UnicastRemoteObject implements DelegatorInterface {
37     private static final long serialVersionUID = 7417123750947132852L;
    private static String host, ftpRootDir;
39     private static int registryPort, ftpServerPort;
    private static LinkedList<File> testSuiteList, testSingleList;
41     private static LinkedList<CustomServerInterface> serverList;

43     public static void main(String[] args) throws RemoteException {
45         // Prevent console logging output
        Logger.getRootLogger().removeAllAppenders();
47         Logger.getRootLogger().addAppender(new NullAppender());

49         // Set IP address to be used by the registry and FTP server on this node
        host = args[0];
51         System.out.println("Using host address: " + host);

53         // Set the port numbers to use for the registry and ftp server (must be
        different)
        registryPort = 12345;
55         ftpServerPort = 12346;

57         // Set the directory the FTP server will use as its home dir
        ftpRootDir = "ftpserver/";

59         // Initialize empty lists for test cases and servers
61         testSuiteList = new LinkedList<File>();
        testSingleList = new LinkedList<File>();
63         serverList = new LinkedList<CustomServerInterface>();

65         // Set the system-wide policy for RMI security (by default, grant all
        permissions)
        System.setProperty("java.security.policy", "rmi.policy");
67

69         // Set the IP address to which a registry will be bound
        System.setProperty("java.rmi.server.hostname", host);

71         // Create the Java RMI registry on this node, using the specified IP address
        and port
        createRegistry();
73    }
```

```

75 // Create the FTP server that will host files
createFTPServer();

77 // Create the lists of test cases that will be sent to servers for execution
createTestList();

79 // Wait for user input before proceeding
81 Scanner scan = new Scanner(System.in);
System.out.println("== Press ENTER when all servers are connected ==");
83 scan.nextLine();

85 // Begin timing here for response time of systme
long start = System.currentTimeMillis();

87 // Update the servers' classpaths and establish connection with the FTP server
89 updateServers();

91 // Run all of the tests by sending them to servers;
// Save them all in this list of results
93 ConcurrentLinkedQueue<Result> results = runTests();

95 // End timing
long end = System.currentTimeMillis();
97 long elapsed = end - start;

99 // Output results from running tests in a formatted way
101 int runs, successes, failures;
runs = successes = failures = 0;

103 for (Result result : results) {
runs += result.getRunCount();
105 failures += result.getFailureCount();

107     for (Failure failure : result.getFailures()) {
System.out.println("EXCEPTION" + failure.getException());
109 System.out.println("TRACE: " + failure.getTrace());
System.out.println("MESSAGE: " + failure.getMessage());
111 System.out.println("HEADER: " + failure.getTestHeader());
System.out.println("DESCRIPTION: " + failure.getDescription());
113     }
}
115 successes = runs - failures;
System.out.println("Total tests run: " + runs);
117 System.out.println("Number of successes: " + successes);
System.out.println("Number of failures: " + failures);
119 System.out.println("Elapsed time: " + elapsed);
}

121 protected Delegator() throws RemoteException {
123     super();
}

125 /**
127  * Create a Java RMI registry for storing a Delegator remote object and any
CustomServer objects that connect with it.
129  * @throws RemoteException
131  */
private static void createRegistry() throws RemoteException {
133     Registry registry = LocateRegistry.createRegistry(registryPort);
Delegator delegator = new Delegator();
registry.rebind("Delegator", delegator);
135 }

137 /**
139  * Create the Apache FTP server the will host files to be used by the
CustomServers.
141  */
private static void createFTPServer() {
143     FtpServerFactory ftpServerFactory = new FtpServerFactory();

// Create one user that will have access to this server.
145 // Currently, the user will be considered an admin and have all permissions
147 UserManager userManager = ftpServerFactory.getUserManager();
BaseUser adminUser = new BaseUser();
adminUser.setName("user");
149 adminUser.setPassword("user");
adminUser.setEnabled(true);
151 ArrayList<Authority> authorities = new ArrayList<Authority>();

```

```

153     authorities.add(new WritePermission());
154     adminUser.setAuthorities(authorities);
155     adminUser.setHomeDirectory(ftpRootDir);
156     adminUser.setMaxIdleTime(999);
157     try {
158         userManager.save(adminUser);
159     } catch (FtpException e2) {
160         e2.printStackTrace();
161     }
162     ftpServerFactory.setUserManager(userManager);
163
164     // Set the listener of the server to the indicated port and host
165     listenerFactory listenerFactory = new ListenerFactory();
166     listenerFactory.setPort(ftpServerPort);
167     listenerFactory.setServerAddress(host);
168     listenerFactory.setIdleTimeout(999);
169
170     // Set some connection properties, such as the number of concurrent logins
171     ConnectionConfigFactory connectionFactory = new ConnectionConfigFactory();
172     connectionFactory.setAnonymousLoginEnabled(true);
173     connectionFactory.setMaxAnonymousLogins(999);
174     connectionFactory.setMaxLogins(999);
175     connectionFactory.setMaxLoginFailures(999);
176     connectionFactory.setMaxThreads(999);
177     ftpServerFactory.setConnectionConfig(connectionFactory.createConnectionConfig());
178
179     // Adjust more connection settings to be as permissive as possible to the user
180     DataConnectionConfigurationFactory dataFactory = new
181     DataConnectionConfigurationFactory();
182     dataFactory.setActiveEnabled(true);
183     dataFactory.setActiveIpCheck(false);
184     dataFactory.setActiveLocalAddress(host);
185     dataFactory.setPassiveAddress(host);
186     dataFactory.setIdleTime(999);
187     listenerFactory.setDataConnectionConfiguration(dataFactory.
188     createDataConnectionConfiguration());
189     ftpServerFactory.addListener("default", listenerFactory.createListener());
190
191     // Allow user direct access to the previously indicated file directory
192     NativeFileSystemFactory fileFactory = new NativeFileSystemFactory();
193     try {
194         fileFactory.createFileSystemView(adminUser);
195         fileFactory.setCreateHome(true);
196     } catch (FtpException e1) {
197         e1.printStackTrace();
198     }
199     ftpServerFactory.setFileSystem(fileFactory);
200
201     // Create the server using these properties, then start it
202     FtpServer ftpServer = ftpServerFactory.createServer();
203
204     try {
205         ftpServer.start();
206     } catch (FtpException e) {
207         e.printStackTrace();
208     }
209 }
210
211 /**
212  * Create list of single tests (testSingleList) and test suites (testSuiteList).
213  * These are obtained from the ftpserver/resources/test-singles and ftpserver/
214  * resources/test-suites directories, respectively.
215  * @throws RemoteException
216  */
217 private static void createTestList() throws RemoteException {
218
219     File file = new File("ftpserver/resources/");
220
221     // Get the nested test-singles and test-suites directories, and
222     // run the appropriate method to parse the given test type
223     for (File subDir : file.listFiles()) {
224         if (subDir.getName().equals("test-singles")) {
225             createSingleTestList(subDir);
226         }
227         else if (subDir.getName().equals("test-suites")) {
228             createTestSuiteList(subDir);
229         }
230     }
231 }

```

```

227     }
228 }
229 /**
230  * Create the testSingleList, which will contain all individual test cases.
231  * This will recursively travel through the parent directory such that all
232  * subdirectories and files are considered.
233  * @param file
234  *       The test_singles directory
235  */
236 private static void createSingleTestList(File file) {
237     System.out.println("begin create single tests");
238     System.out.println(FilenameUtils.getExtension(file.getName()));
239
240     // If file is a directory, then recursively call this method again
241     if (file.isDirectory()) {
242         for (File subDir : file.listFiles()) {
243             createSingleTestList(subDir);
244         }
245     }
246     // Only consider files in the directory that end in .class
247     else if (FilenameUtils.getExtension(file.getName()).equals("class") && (!file.
248         getName().contains("$"))) {
249         testSingleList.add(file);
250     }
251     System.out.println("end create single tests");
252 }
253
254 /**
255  * Create the testSuiteList, which will contain all test suites.
256  * This will recursively travel through the parent directory such that all
257  * subdirectories and files are considered.
258  * @param file
259  *       The test_suites directory
260  */
261 private static void createTestSuiteList(File file) {
262     System.out.println("begin create suite tests");
263
264     // If file is a directory, then recursively call this method again
265     if (file.isDirectory()) {
266         for (File subDir : file.listFiles()) {
267             createTestSuiteList(subDir);
268         }
269     }
270     // Only consider files in the directory that end in .class
271     else if (FilenameUtils.getExtension(file.getName()).equals("class")) {
272         testSuiteList.add(file);
273     }
274     System.out.println("end create suite tests");
275 }
276
277 /**
278  * This method is called by a CustomServer in order to bind that server to the
279  * registry on this node.
280  * @param remoteObject
281  *       The reference to the remote CustomServer
282  */
283 public void rebindServer(CustomServerInterface remoteObject) throws
284     RemoteException {
285     System.out.println("Starting to bind server");
286     try {
287         Registry registry = LocateRegistry.getRegistry(host, registryPort);
288         registry.rebind("CustomServer-" + registry.list().length, remoteObject);
289
290         // Add this server to the server list for easy access later
291         serverList.add(remoteObject);
292
293         // Announce the current list of objects in the registry
294         System.out.println("CURRENTLY BOUND REMOTE OBJECTS: ");
295         for (String s : registry.list()) {
296             System.out.println(s);
297         }
298     } catch (Exception e) {
299         e.printStackTrace();
300     }
301 }

```

```

301 }
303 /**
304  * Call each of the servers to update their classpaths and retrieve any
305  * necessary configuration files needed to run the tests
306  *
307  * @throws RemoteException
308  */
309 private static void updateServers() throws RemoteException {
310     try {
311         System.out.println("Begin update servers");
312
313         for (CustomServerInterface server : serverList) {
314             server.updateClassLoader();
315         }
316         System.out.println("End update servers");
317     } catch (Exception e) {
318         e.printStackTrace();
319     }
320 }
321
322 /**
323  * This controls all of the load sharing used for distributing tests to servers.
324  *
325  * @return The list of Results obtained from running tests on the servers
326  * @throws RemoteException
327  */
328 private static ConcurrentLinkedQueue<Result> runTests() throws RemoteException {
329     try {
330         // Use ConcurrentLinkedQueues to allow modification from multiple threads
331         ConcurrentLinkedQueue<Result> resultQueue = new ConcurrentLinkedQueue<Result>();
332         ConcurrentLinkedQueue<TestAgent> agents = new ConcurrentLinkedQueue<TestAgent>();
333
334         // Add separate agent for each server
335         for (CustomServerInterface server : serverList) {
336             agents.add(new TestAgent(server, resultQueue));
337         }
338
339         // Class loader for converting bytes to Class object (and nothing else)
340         SimpleClassLoader simpleLoader = new SimpleClassLoader();
341
342         // Iterators for each test list
343         Iterator<File> suiteIterator = testSuiteList.iterator();
344         Iterator<File> singleIterator = testSingleList.iterator();
345
346         // Iterate first through the list of single test cases
347         while (singleIterator.hasNext()) {
348             // Access (don't remove) the first test in the list
349             File testFile = testSingleList.getFirst();
350
351             // Remove the first agent in the queue
352             TestAgent agent = agents.remove();
353
354             // Consider agents that are not alive (ie not running tests)
355             if (!agent.isAlive()) {
356                 // Point agent to new agent with the same server reference, but with the
357                 // current test
358                 agent = new TestAgent(agent, testFile);
359
360                 // Have the agent execute the server's runTest method in its own thread
361                 agent.start();
362
363                 // Remove the test from the list
364                 singleIterator.next();
365             }
366
367             // Add the agent back to the queue
368             agents.add(agent);
369         }
370
371         // Iterate next through all test suites
372         while (suiteIterator.hasNext()) {
373             // Convert the File from bytes, then to a Class

```



```

377     FileInputStream in = new FileInputStream(suiteIterator.next());
378     byte[] classBytes = new byte[(int) testSuiteList.getFirst().length()];
379     in.read(classBytes);
380     simpleLoader = new SimpleClassLoader();
381     Class<?> convertedClass = simpleLoader.convertToClass(classBytes);
382
383     // Get all test cases (classes) referenced in the test suite
384     SuiteClasses suiteAnnotation = convertedClass.getAnnotation(SuiteClasses.
class);
385     Class<?>[] classesInSuite = suiteAnnotation.value();
386
387     // Iterate through all test cases in the test suite
388     for (int i = 0; i < classesInSuite.length; i++) {
389         // Access (don't remove) the first test in the list
390         Class<?> c = classesInSuite[i];
391
392         // Remove the first agent in the queue
393         TestAgent agent = agents.remove();
394
395         // Consider agents that are not alive (ie not running tests)
396         if (!agent.isAlive()) {
397             System.out.println("Assigning test: " + c.getName());
398
399             // Point agent to new agent with the same server reference, but with
the current test
400             agent = new TestAgent(agent, c);
401
402             // Have the agent execute the server's runTest method in its own
thread
403             agent.start();
404
405             // Move to the next test case in the list
406             i++;
407         }
408
409         // Add the agent back to the queue
410         agents.add(agent);
411     }
412
413     // Wait until all agents have finished before returning result
414     while (!agents.isEmpty()) {
415         TestAgent agent = agents.peek();
416         if (!agent.isAlive()) {
417             agents.remove();
418         }
419     }
420
421     return resultQueue;
422 } catch (Exception e) {
423     e.printStackTrace();
424     return null;
425 }
426 }
427
428 /** Debugging method to show remote communication between nodes */
429 public String ping() throws RemoteException {
430     return "==PING!==" ;
431 }
432 }

```

../code/src/delegator/Delegator.java

7.3.2 CustomServer

```
package server;

import java.io.File;
import java.io.FileInputStream;
import java.io.FileOutputStream;
import java.io.IOException;
import java.net.MalformedURLException;
import java.net.URL;
import java.rmi.Naming;
import java.rmi.NotBoundException;
import java.rmi.RemoteException;
import java.rmi.server.UnicastRemoteObject;

import org.apache.commons.net.ftp.FTPClient;
import org.apache.commons.net.ftp.FTPFile;
import org.junit.runner.JUnitCore;
import org.junit.runner.Result;

import delegator.DelegatorInterface;
import delegator.SimpleClassLoader;

public class CustomServer extends UnicastRemoteObject implements
    CustomServerInterface {

    private static final long serialVersionUID = -3247691930817118343L;
    private static String host, ftpClassDir, ftpJarDir, ftpRootDir, userName,
        userPassword;
    private static int registryPort, ftpServerPort;
    private static FTPClient client;

    public static void main(String[] args) throws RemoteException,
        MalformedURLException, NotBoundException {

        // Immediately set the default class path to a new CustomClassLoader
        System.setProperty("java.system.class.loader", "server.CustomClassLoader");

        // Get the IP address for the delegator from the command line
        host = args[0];

        // Some nodes may need to assign their own rmi server hostname with their
        // ip address to allow communication with Delegator
        if (args.length > 1) {
            String thisRMIhost = args[1];
            System.setProperty("java.rmi.server.hostname", thisRMIhost);
        }

        // Initialize various vars
        System.out.println("Using host address: " + host);
        registryPort = 12345;
        ftpServerPort = 12346;
        ftpClassDir = "resources/bin/";
        ftpJarDir = "resources/lib/";
        ftpRootDir = "resources/";
        userName = "user";
        userPassword = "user";

        // Set a very permissive RMI policy
        System.setProperty("java.security.policy", "rmi.policy");

        // Retrieve a remote reference to the Delegator stored in the registry
        DelegatorInterface delegator = (DelegatorInterface) Naming.lookup("//" + host
            + ":" + registryPort + "/Delegator");

        // PING the delegator to ensure patent communication
        System.out.println(delegator.ping());

        // Bind an instance of a CustomServer to the Delegator's registry
        CustomServer server = new CustomServer();
        delegator.rebindServer((CustomServerInterface) server);
    }

    protected CustomServer() throws RemoteException {
        super();
    }

    /** Update this server's CustomClassLoader to points to the files on the
        Delegator,
        * accessed through its FTP server
    */
}
```

```

74  */
75  public void updateClassLoader() {
76
77      System.out.println("BEGIN UPDATECLASSLOADER");
78      try {
79          // Use the Apache Commons Net library to facilitate FTP communication
80          client = new FTPClient();
81          client.connect(host, ftpServerPort);
82          client.login(userName, userPassword);
83          client.setKeepAlive(true);
84
85          // Set classpath to the class directory for the resources on the Delegator
86          CustomClassLoader.addURLToSystemClassLoader(new URL("ftp://" + userName + ":"
87 + userPassword + "@" + host + ":" + ftpServerPort + "/" + ftpClassDir));
88
89          // Set classpath to each jar file for the resources on the Delegator
90          FTPFile[] jarFiles = client.listFiles(ftpJarDir);
91          for (int i = 0; i < jarFiles.length; i++) {
92              CustomClassLoader.addURLToSystemClassLoader(new URL("ftp://" + userName +
93 + userPassword + "@" + host + ":" + ftpServerPort + "/" + ftpJarDir +
94 + jarFiles[i].getName()));
95          }
96
97          // Add any supplemental configuration files needed to run subsequent tests
98          // (located immediately in the ftpserver/ directory on the Delegator node)
99          FTPFile[] files = client.listFiles();
100         for (int i = 0; i < files.length; i++) {
101
102             FTPFile file = files[i];
103             if (!file.getName().equals("resources")) {
104                 createResources(file, "");
105             }
106         }
107     } catch (Exception e) {
108         e.printStackTrace();
109     }
110
111     System.out.println("END UPDATECLASSLOADER");
112 }
113
114 /** This method will recursively go through the home directory from the FTP
115     server
116     * and retrieve and recreate the directory structure for all files (NOT in the
117     main
118     * resources directory; ie only supplemental, non-class files needed to execute
119     the
120     * given test cases)
121     */
122 private void createResources(FTPFile file, String parentDir) {
123     try {
124         if (file.isDirectory()) {
125             File targetFile = new File(parentDir + file.getName() + "/");
126             targetFile.mkdir();
127             for (FTPFile subDir : client.listFiles(file.getName())) {
128                 createResources(subDir, parentDir + file.getName() + "/" + subDir.getName());
129             }
130         } else {
131             client.retrieveFile(parentDir + file.getName(), new FileOutputStream(
132                 parentDir + file.getName()));
133         }
134     } catch (IOException e) {
135         e.printStackTrace();
136     }
137 }
138
139 /** Execute the test case (formatted as a Class object) retrieved from the
140     Delegator */
141 public Result runTest(Class testClass) {
142
143     System.out.println("**Running test: " + testClass.getName());
144     try {
145         // Run the test case using JUnit, storing and sending the result back to the
146         Delegator
147         Result result = JUnitCore.runClasses(testClass);
148
149         return result;
150     } catch (Exception e) {

```

```

144         e.printStackTrace();
145         return null;
146     }
147 }
148
149 /** Execute the test case (formatted as a File object) retrieved from the
150     Delegator */
151 public Result runTest(File testFile) {
152     System.out.println("**Running test: " + testFile.getName());
153
154     try {
155         // Convert the File to bytes, and then to a Class object
156         FileInputStream in = new FileInputStream(testFile);
157         byte[] classBytes = new byte[(int) testFile.length()];
158         in.read(classBytes);
159         SimpleClassLoader loader = new SimpleClassLoader();
160         Class convertedClass = loader.convertToClass(classBytes);
161
162         // Run the test case using JUnit, storing and sending the result back to the
163         Delegator
164         Result result = JUnitCore.runClasses(convertedClass);
165
166         return result;
167     } catch (Exception e) {
168         e.printStackTrace();
169         return null;
170     }
171 }
172
173 /** Debugging method to ensure patent communication between server and delegator
174     */
175 public String ping() throws RemoteException {
176     return "=====PING FROM SERVER=====";
177 }
178 }

```

../code/src/server/CustomServer.java

7.3.3 CustomClassLoader

```
1 package server;
3 import java.beans.IntrospectionException;
import java.io.File;
5 import java.io.IOException;
import java.lang.reflect.Method;
7 import java.net.MalformedURLException;
import java.net.URL;
9 import java.net.URLClassLoader;

11 public class CustomClassLoader extends URLClassLoader {
13     /** This constructor is used when this class loader is initially set as the
        * default system
        * class loader. It will set the initial classpath using URLs obtained from
        * getNestedURLs
        */
15     public CustomClassLoader(ClassLoader l) throws IOException {
17         super(getNestedURLs(), l.getParent());
19     }

21     /** This method will return URLs for all jar files in the local lib directory to
        * the classpath,
        * as well as the URL for the local bin directory. This is used for help with
        * initialization
        */
23     private static URL[] getNestedURLs() {
25         String binDirString = CustomClassLoader.class.getProtectionDomain().
            getCodeSource().getLocation().getPath();
            String libDirString = binDirString.substring(0, binDirString.length() - 4) + "
            lib/";

27         File libDirFile = new File(libDirString);
29         File[] jarFiles = libDirFile.listFiles();
            URL[] urls = new URL[jarFiles.length + 1];
31         try {
33             urls[0] = new URL("file://" + binDirString);

35             for (int i = 1; i < urls.length; i++) {
37                 urls[i] = new URL(jarFiles[i - 1].toURI().toString());
39             } catch (MalformedURLException e) {
41                 e.printStackTrace();
43             }

45             return urls;
47         }

49     /** Add a specific URL to the class path.
        * In order to facilitate this method, reflection needed to be use to access the
        * method
        * since it is protected. The following technique was obtained from
        * http://baptiste-wicht.com/posts/2010/05/tip-add-resources-dynamically-to-a-
        * classloader.html
        */
51     public static void addURLToSystemClassLoader(URL url) throws
        IntrospectionException {
53         URLClassLoader systemClassLoader = (URLClassLoader) ClassLoader.
            getSystemClassLoader();
            Class<URLClassLoader> classLoaderClass = URLClassLoader.class;

55         try {
57             Method method = classLoaderClass.getDeclaredMethod("addURL", new Class[] {
                URL.class });
                method.setAccessible(true);
                method.invoke(systemClassLoader, new Object[] { url });
59             } catch (Throwable t) {
61                 t.printStackTrace();
                throw new IntrospectionException("Error when adding url to system
                ClassLoader ");
            }

63     /** Deprecated method: Trying to replicate the findClass() method of
        URLClassLoader using reflection*/
```

```

65 public static Class findClassWithSystemClassLoader(String name) throws
    IntrospectionException {
66     URLClassLoader systemClassLoader = (URLClassLoader) ClassLoader.
        getSystemClassLoader();
67     Class<URLClassLoader> classLoaderClass = URLClassLoader.class;
68
69     Class convertedClass = null;
70
71     for (Method m : classLoaderClass.getDeclaredMethods()) {
72         System.out.print(m.getName() + " ");
73
74         for (Class c : m.getParameterTypes()) {
75             System.out.print(c.getName() + " ");
76         }
77
78         System.out.println(", return: " + m.getReturnType().getName());
79     }
80     try {
81         Method method = classLoaderClass.getDeclaredMethod("findClass", new Class[]
            { String.class });
82         method.setAccessible(true);
83         convertedClass = (Class) method.invoke(systemClassLoader, new Object[] {
            name });
84     } catch (Throwable t) {
85         t.printStackTrace();
86         throw new IntrospectionException("Error when finding class with system
            ClassLoader ");
87     }
88     return convertedClass;
89 }

```

../code/src/server/CustomClassLoader.java

7.3.4 TestAgent

```
package delegator;

2
import java.io.File;
4
import java.rmi.RemoteException;
import java.util.concurrent.ConcurrentLinkedQueue;
6
import org.junit.runner.Result;
8
import server.CustomServerInterface;
10
public class TestAgent extends Thread implements Runnable {
12
    private Class<?> testClass;
14
    private static ConcurrentLinkedQueue<Result> resultQueue;
    private CustomServerInterface server;
16
    private File testFile;

18
    /**
     * Constructor for assigning the server used by this agent and the shared result
     * queue that all results will be added to.
20
     */
    public TestAgent(CustomServerInterface server, ConcurrentLinkedQueue<Result>
        resultQueueRef) {
22
        this.server = server;
        if (resultQueue == null) {
24
            resultQueue = resultQueueRef;
        }
26
    }

28
    /**
     * Constructor for assigning an existing agent a new test class.
     * This constructor is needed, since thread will "die" and not be usable after
     * executing run() method once.
30
     */
    public TestAgent(TestAgent agent, Class<?> testClass) {
32
        server = agent.getServer();
        resultQueue = agent.getResultQueue();
34
        this.testClass = testClass;
    }
36

38
    /**
     * Constructor for assigning an existing agent a new test file.
     * This constructor is needed, since thread will "die" and not be usable after
     * executing run() method once.
40
     */
    public TestAgent(TestAgent agent, File testFile) {
42
        server = agent.getServer();
        resultQueue = agent.getResultQueue();
44
        this.testFile = testFile;
    }
46

48
    /**
     * The run method is executed whenever the super method from Thread, start(), is
     * executed.
     * It will run the appropriate runTest() method for the CustomServer, depending
     * on whether it has a test file or test class assigned
50
     */
    @Override
52
    public void run() {
54
        try {
            if (testFile == null) {
56
                Result result = server.runTest(testClass);
                addResult(result);
58
            }
            else {
60
                Result result = server.runTest(testFile);
                addResult(result);
62
            }
        } catch (RemoteException e) {
64
            e.printStackTrace();
        }
66
    }

68
    /**
     * Synchronized method to allow multiple, different agents to add to the same
     * result queue.
70
     */
    public static synchronized void addResult(Result result) {
```

```
72     resultQueue.add(result);
73 }
74 private ConcurrentLinkedQueue<Result> getResultQueue() {
75     return resultQueue;
76 }
77
78 private CustomServerInterface getServer() {
80     return server;
81 }
82
83 public void setTest(Class<?> testClass) {
84     this.testClass = testClass;
85 }
86 }
```

../code/src/delegator/TestAgent.java

References

- [1] Avishay Traeger, Erez Zadok, Nikolai Joukov, and Charles P Wright. A nine year study of file system and storage benchmarking. *ACM Transactions on Storage (TOS)*, 4(2):5, 2008.